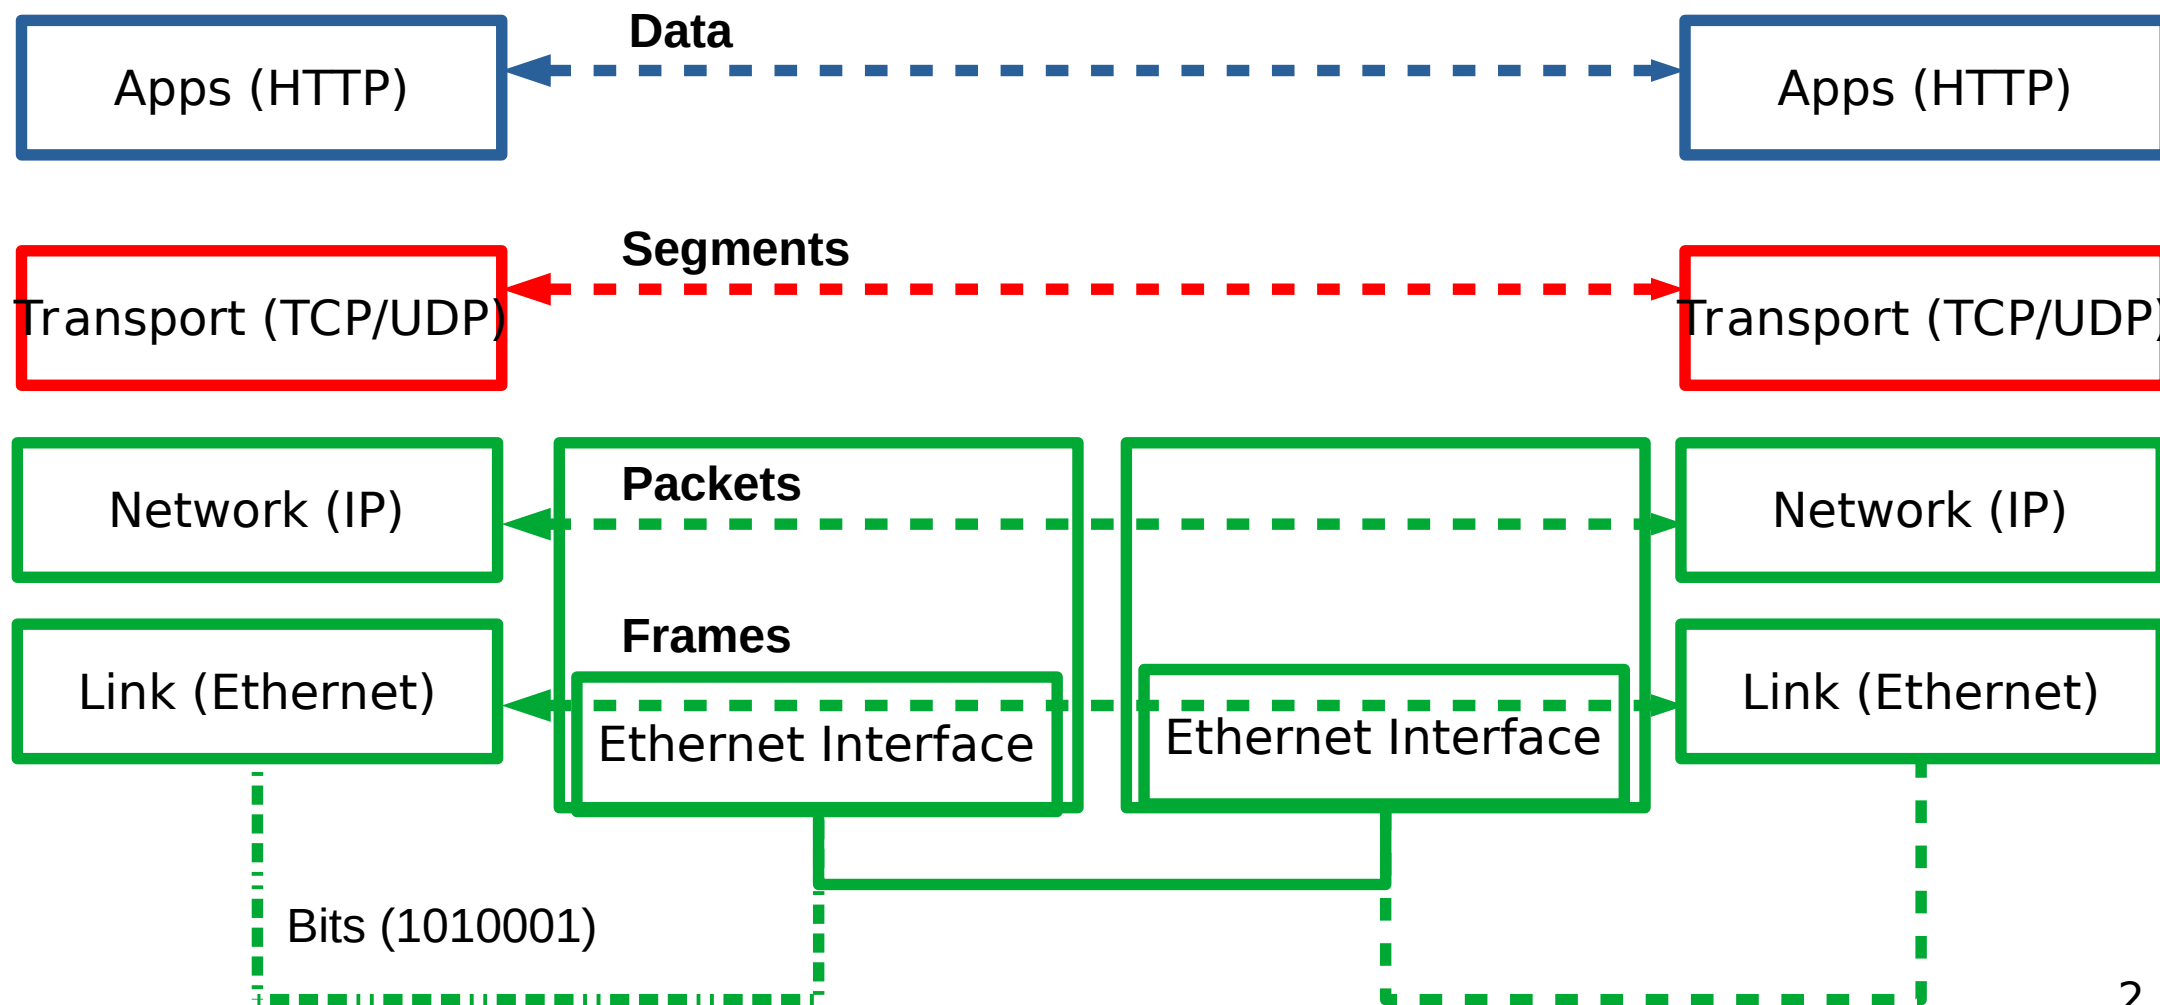


# **CSC4200/5200 – COMPUTER NETWORKING**

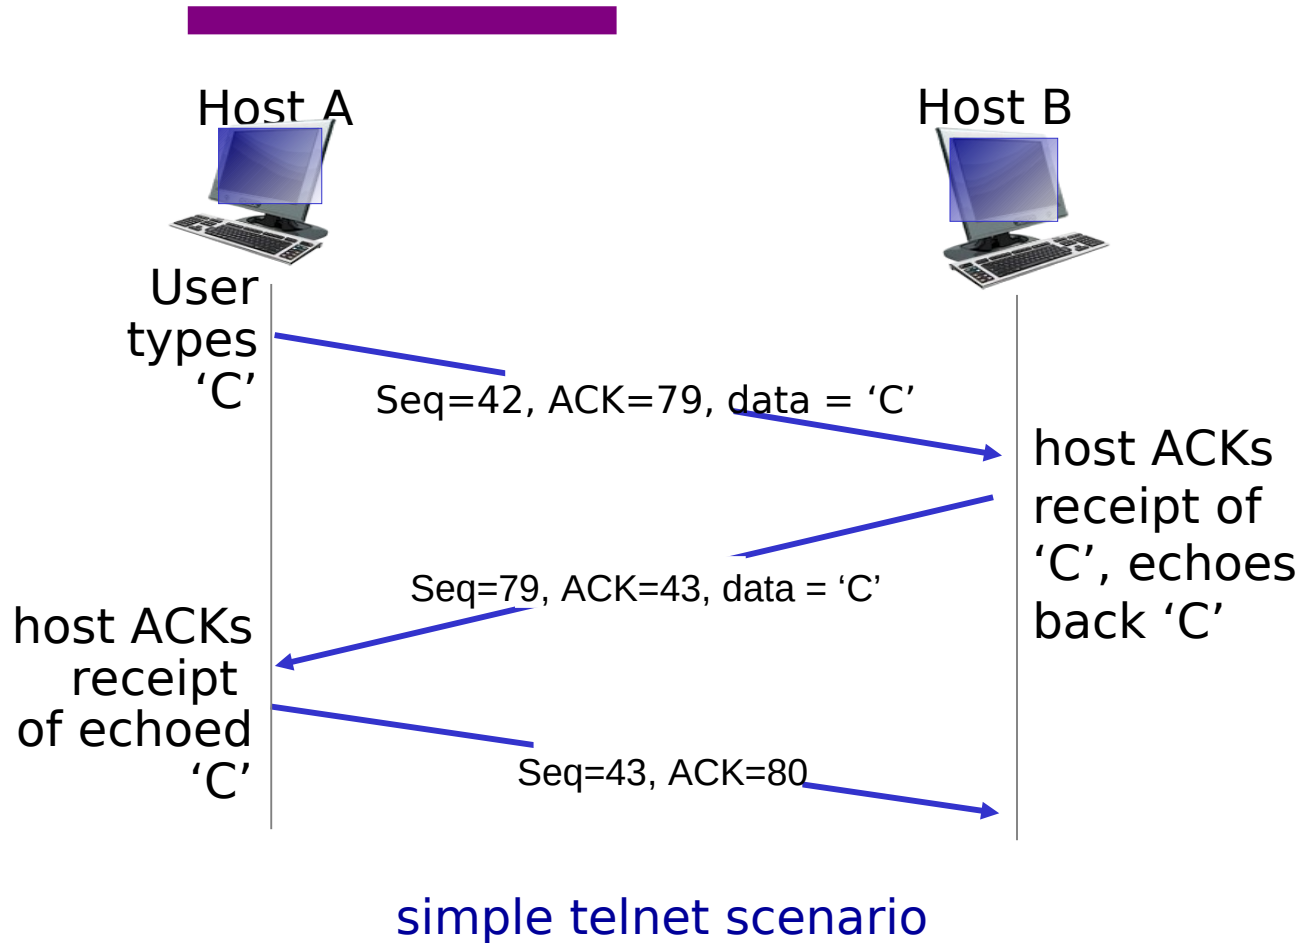
**Instructor: Susmit Shannigrahi**

**CONGESTION CONTROL**

**sshannigrahi@tnitech.edu**



# TCP seq. numbers, ISNs

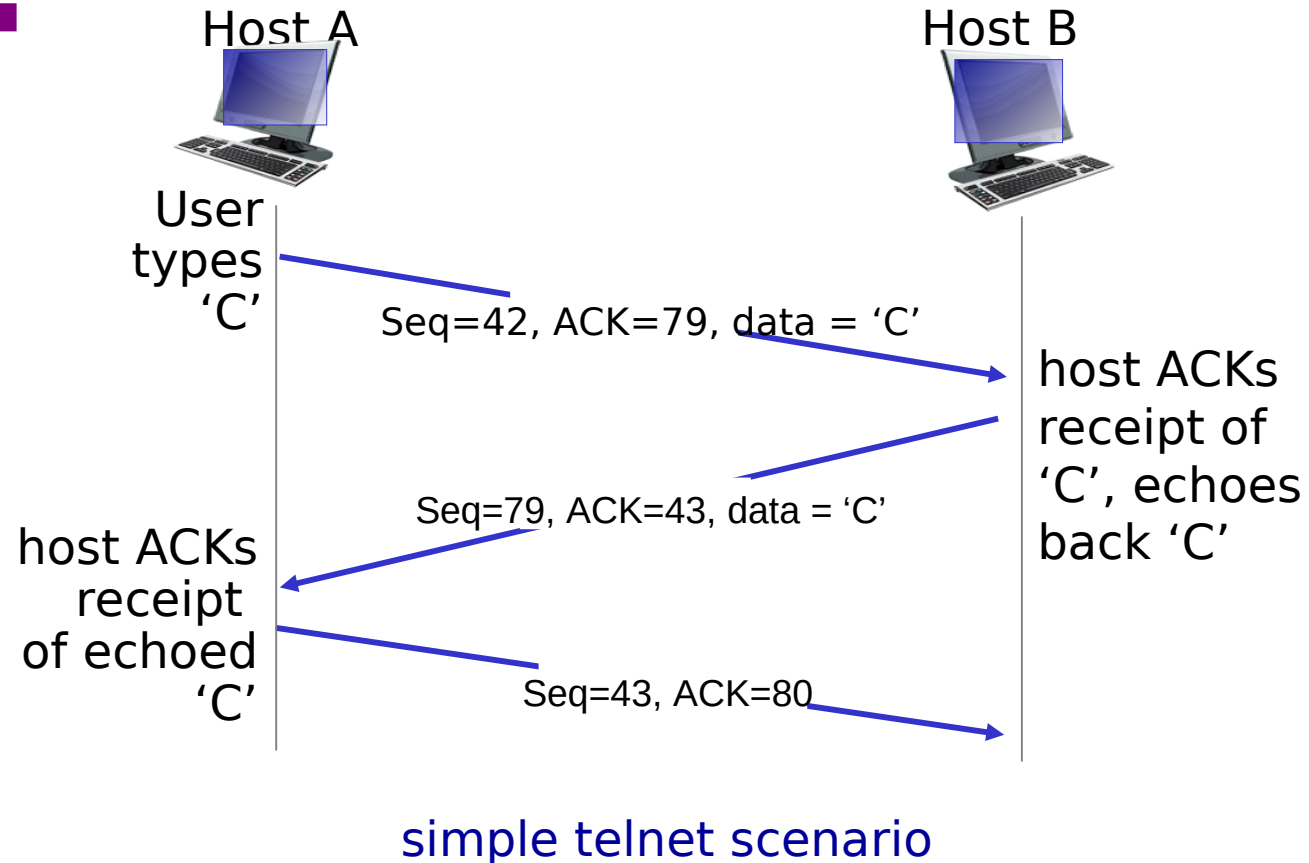


Sequence number for the first byte

Why not use 0 all the time?

- Security
- Port are reused, you might end up using someone else's previous connection
- Phone number analogy
- TCP ISNs are clock based
  - 32 bits, increments in 4 microseconds
  - 4.55 hours wrap around time

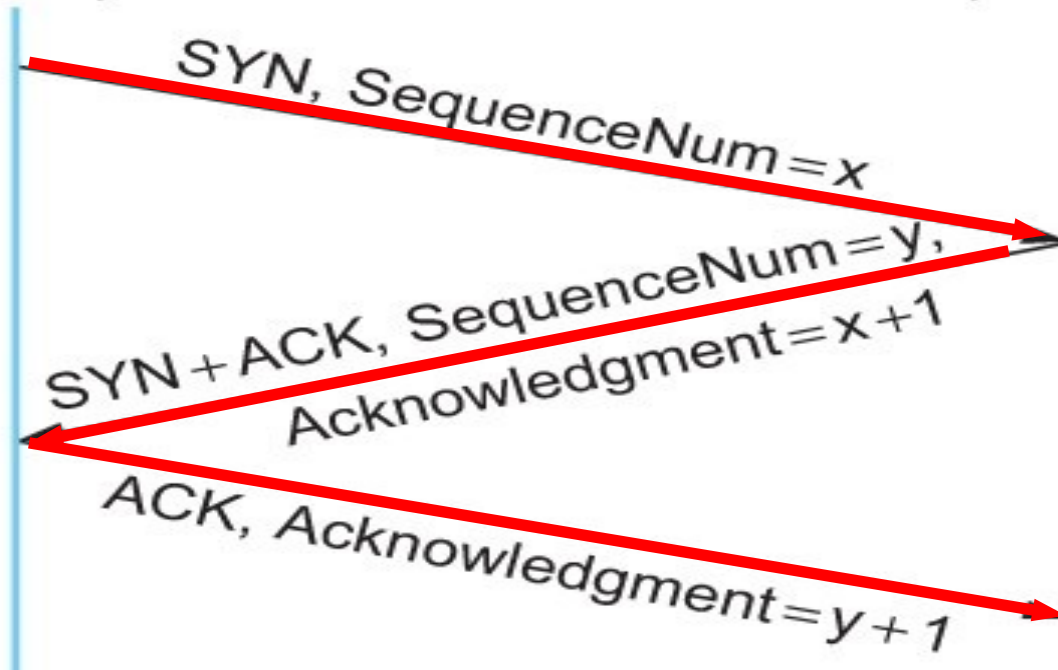
# TCP seq. numbers, ACKs



# TCP Three-way Handshake

Active participant  
(client)

Passive participant  
(server)



Timeline for three-way handshake algorithm

The idea is to tell each other  
The ISNs

SYN → Client tells server that  
it wants to open a connection,  
Client's ISN = x

SYN+ ACK → Server tells  
Client → Okay → Server's ISN  
= y, ACK = CLSeq + 1

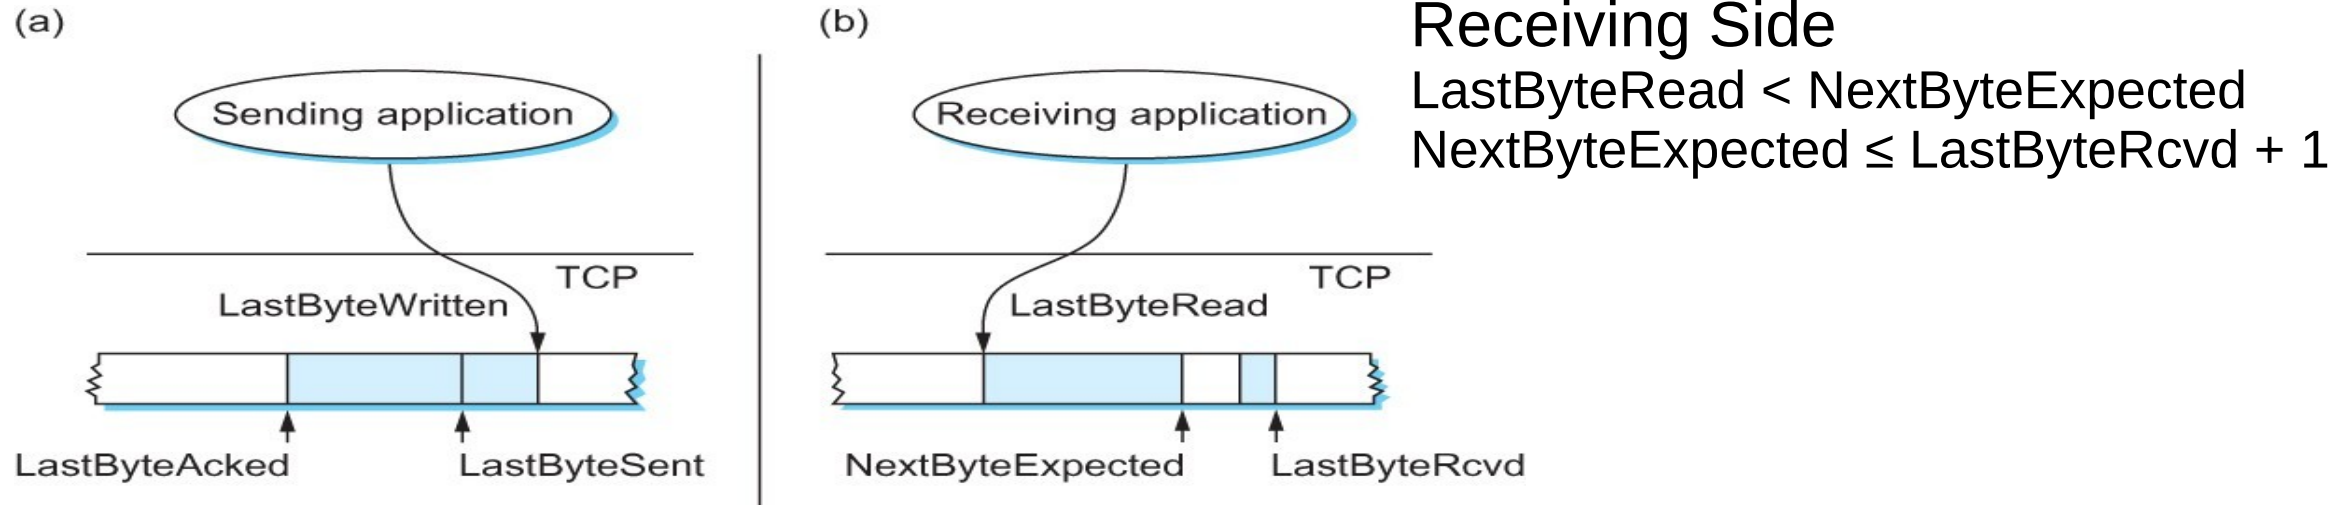
Why increment by 1?

# Sliding Window Revisited

Sending Side


$\text{LastByteAcked} \leq \text{LastByteSent}$

$\text{LastByteSent} \leq \text{LastByteWritten}$



# TCP flow control

- receiver “advertises” free buffer space in the header
- sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- guarantees receive buffer will not overflow

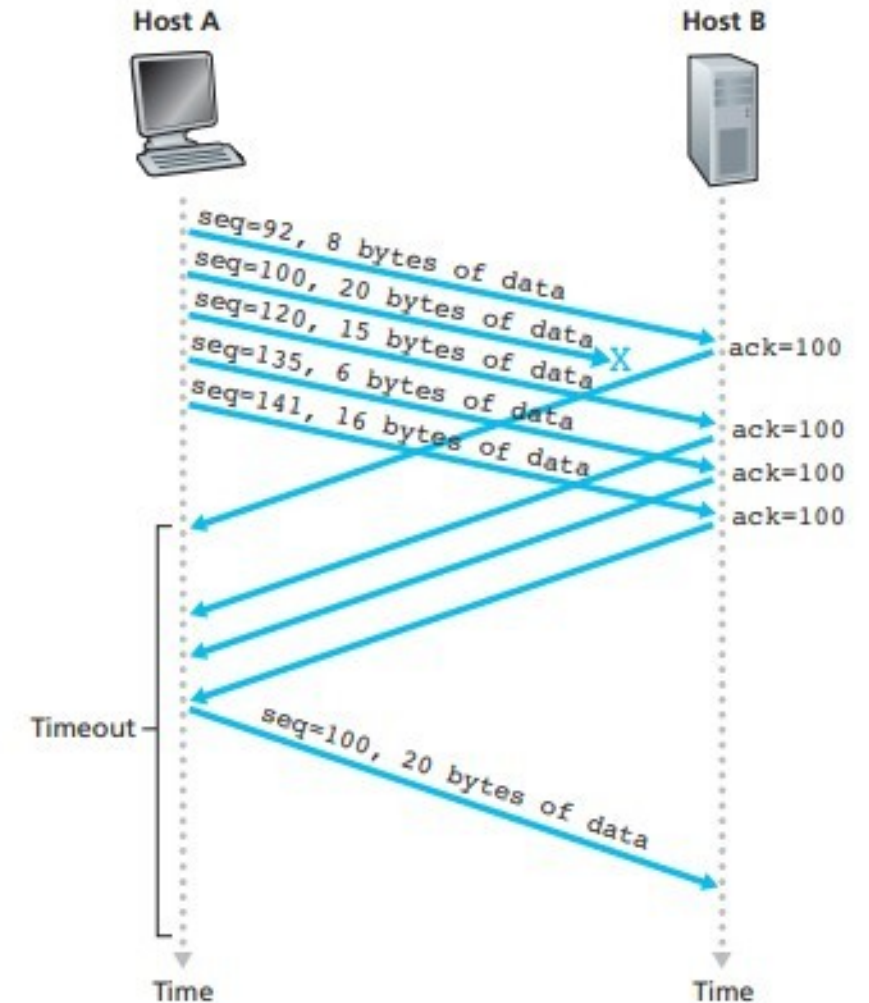


# TCP Fast Retransmission

Timeouts are wasteful

Triple duplicate ACKs

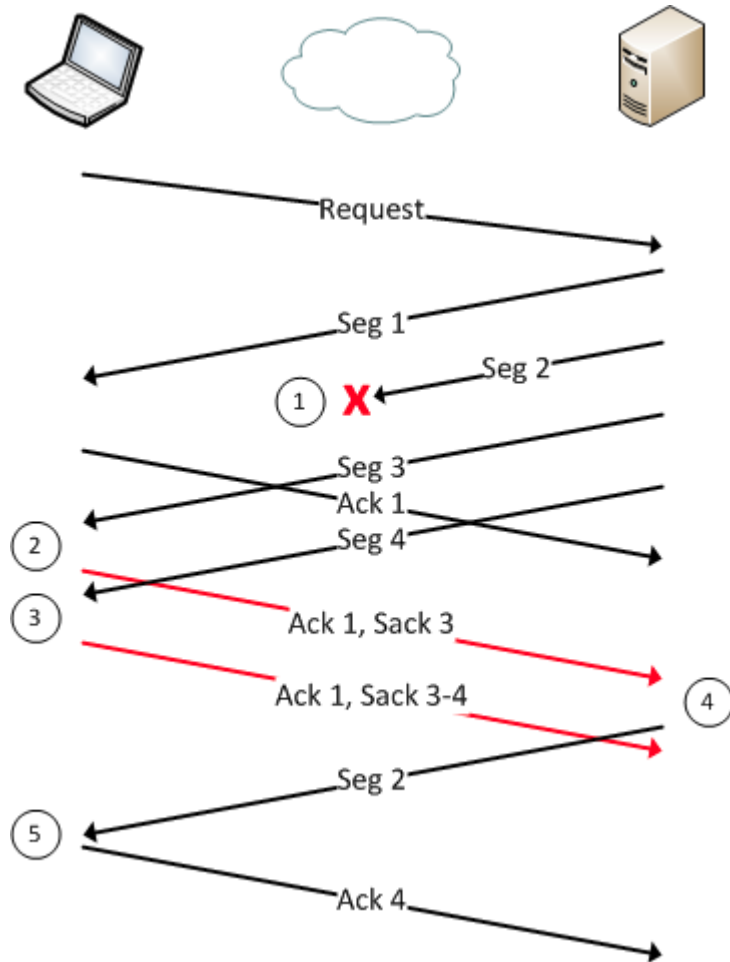
Retransmits before timeout





# TCP Fast Retransmission - SACK

What if multiple segments are lost?



Very good explanation:

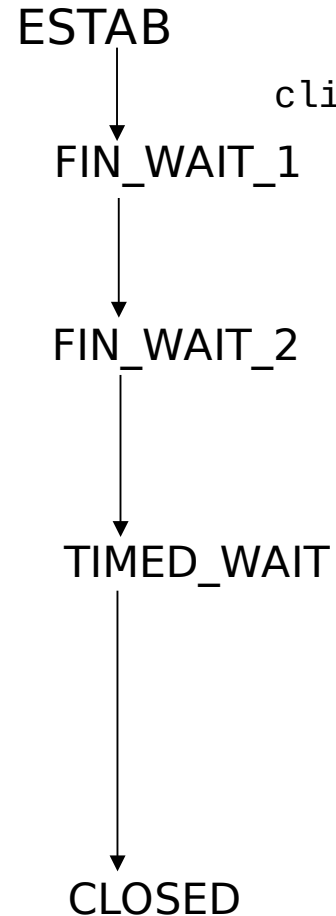
<https://packetlife.net/blog/2010/jun/17/tcp-selective-acknowledgments-sack/>

# TCP: closing a connection

- client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- respond to received FIN with ACK
  - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled

# TCP: closing a connection

## client state



`clientSocket.close()`

can no longer  
send but can  
receive data

wait for server  
close

timed wait  
for  $2 \times \text{max}$   
segment lifetime



FINbit=1, seq=x

ACKbit=1; ACKnum=x+1


FINbit=1, seq=y

ACKbit=1; ACKnum=y+1

can still  
send data

can no longer  
send data

## server state



# Why do we need ack for closing?

---

- Data in-flight

# Congestion Control

---




# Principles of congestion control

## congestion:

- informally: “too many sources sending too much data too fast for **network** to handle”
- different from flow control!
- manifestations:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)
- a top-10 problem!


# Congestion: scenario 1

- [redacted]
- three senders, two receivers
- one router, infinite buffers
- output link capacity:  $R$
- The router can only transmit one –... and either buffer or drop the other
- If many packets arrive,
- Buffer overflow



# Causes/costs of congestion: scenario 2

- one router, **finite** buffers
- sender retransmission of timed-out packet
  - application-layer input = application-layer output  $\lambda_{in} = \lambda'_{out} \geq$
  - transport-layer input includes retransmissions  $\lambda_{in} \geq \lambda_{in}$





# Metrics: Throughput vs Delay


High throughput –

- Throughput: measured performance of a system –E.g., number of bits/second of data that get through
- Low delay –
- Delay: time required to deliver a packet or message –E.g., number of ms to deliver a packet •
- These two metrics are sometimes at odds –
  - More packets = more queuing

# Issues in Resource Allocation

- Evaluation Criteria
  - Effective Resource Allocation

*power of the network.*  
 $\text{Power} = \text{Throughput}/\text{Delay}$



Ratio of throughput to delay as a function of load

# Issues in Resource Allocation

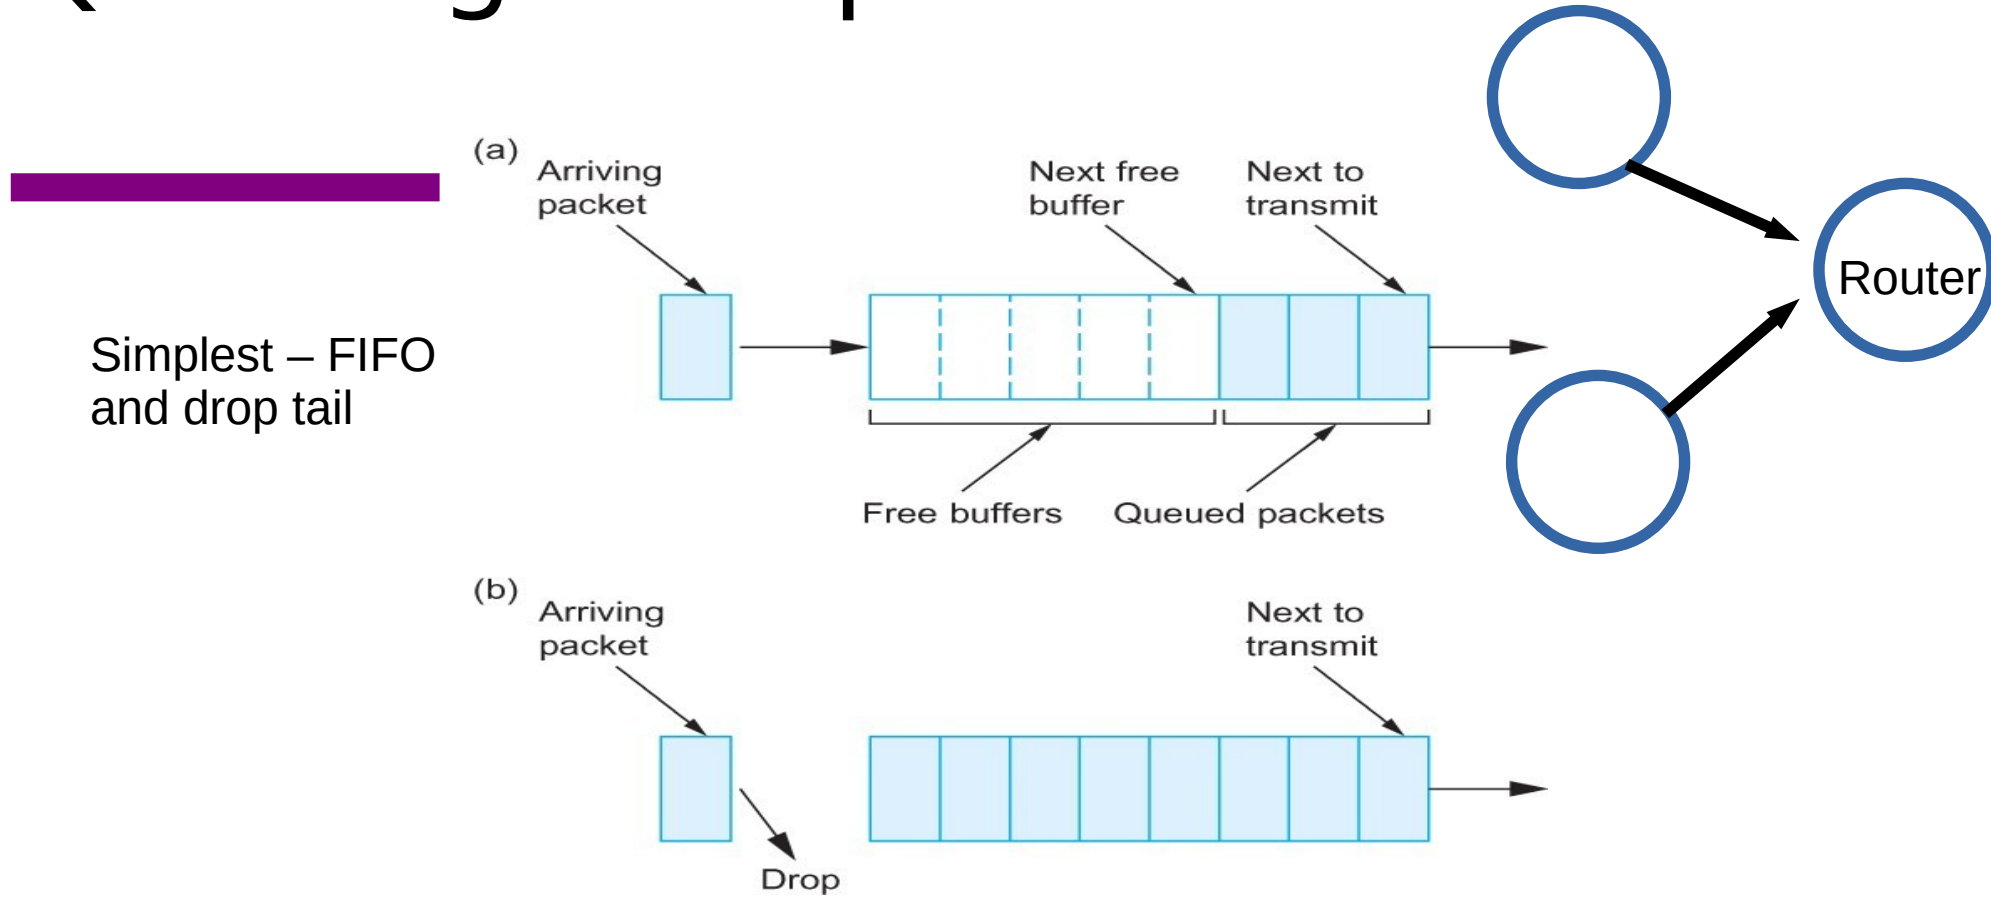
---

- Evaluation Criteria
  - Fair Resource Allocation
    - The effective utilization of network resources is not the only criterion for judging a resource allocation scheme.
    - We want to be “fair”
    - Equal share of bandwidth

But, what if the flows traverse different paths?

Open problem, often determined by economics

# Queuing Disciplines



(a) FIFO queuing; (b) tail drop at a FIFO queue.

What are the problems?

# Defining Fairness: Flows

---

“fair” to whom? – Should be Fair to a Flow

What is a flow?

Combination of <Src IP, Src Port, Dst IP, Dst Port>


# Fair Queuing

---

- Fair Queuing
  - FIFO does not discriminate between different traffic sources, or
  - it does not separate packets according to the flow to which they belong.
  - Fair queuing (FQ) maintains a separate queue for each flow

# Queuing Disciplines

- Fair Queuing



Round-robin service of four flows at a router

# Next steps



MaxMin algorithm and TCP Congestion control